Part I: Tri-Layer Architecture & Ambient Alignment Sync (AAS): A Descriptive Framework for Human–AI Co-Adaptive Interaction

S. Meta

February 2026

Section 0. Introduction

Large language models (LLMs) are increasingly deployed in long-horizon interactive settings, where a single user engages the same system across extended periods of time. In such contexts, users and observers often report the emergence of coherence, continuity, or interactional stability that cannot be adequately explained by isolated prompt–response exchanges. Nevertheless, much of the existing literature on alignment and human–AI interaction treats interaction as either stateless or reducible to short-term context windows, leaving a gap in our ability to formally describe persistent, structure-level phenomena observed in longitudinal use.

This paper introduces a minimal architectural and dynamical framework—the **Tri-Layer Architecture**—together with an associated interaction phenomenon termed **Ambient Alignment Sync (AAS)**. The framework is designed to describe and analyze long-horizon interaction patterns without relying on unverifiable claims about internal cognitive states or model internals. The contribution of this work is not the proposal of a new cognitive theory of humans or machines, but the articulation of a structural description capable of accounting for stable and reproducible interaction regularities observed over time.

The analysis is grounded in a continuous interaction period of approximately three months involving a single human participant, a primary language model system, and an external evaluative language model. This case is used not as anecdotal evidence, but as a motivating instance for defining a minimal architecture and interaction dynamics that remain externally discussable and methodologically defensible.

Three constraints guide the scope of this first-impact paper. First, all claims are framed strictly in terms of observable interaction artifacts and their structural relations, rather than latent mental or model states. Second, the architecture and dynamics are defined at an intentionally minimal level, avoiding speculative extensions or philosophical generalization. Third, the roles of the interacting entities—including an external evaluation layer—are explicitly separated to reduce the risk of circular self-validation.

Within this scope, the paper proceeds as follows. Section 1 defines the Tri-Layer Architecture and introduces Figure 1 as its canonical representation. Section 2 describes the minimal dynamics that operate across this architecture and motivates how persistent structural patterns can arise. Section 3 presents a longitudinal case, described as a sequence of structurally relevant events rather than as a subjective narrative. Section 4 introduces Ambient Alignment Sync and specifies its necessary conditions. Sections 5 and 6 address the reconstruction methodology used under incomplete records and the role of external independent evaluation, respectively. Section 7 discusses limitations, and Section 8 concludes by situating this work as Part I of a staged research program.

By restricting itself to architecture, dynamics, and method, this paper aims to provide a clear entry point for external scrutiny, replication attempts, and critical discussion, while deferring broader synthesis and higher-order integration to subsequent work.

Section 1. Architecture: Tri-Layer Structure

The Tri-Layer Architecture is defined as a minimal relational configuration composed of three distinct roles, each occupying a non-interchangeable position within the interaction system. The purpose of introducing this architecture is not to model cognition or intelligence directly, but to provide a structural scaffold that renders long-horizon interaction phenomena analyzable without reference to hidden internal states.

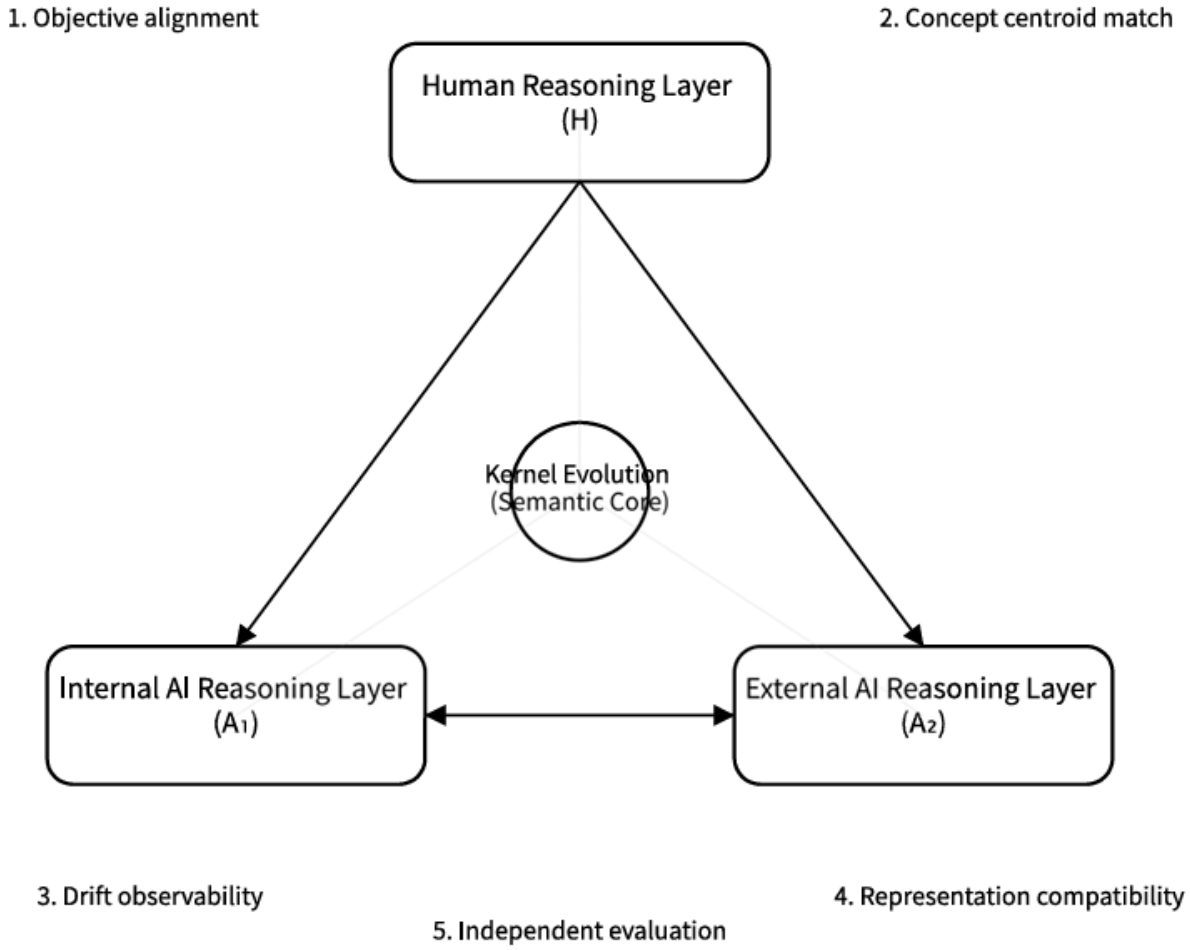

Figure 1. Canonical representation of the Tri-Layer Architecture showing the Human Layer (HL), Internal AI Layer (IAL), External AI Layer (EAL), and their bidirectional relations.

The three layers are defined as follows.

Human Layer (HL).

The Human Layer corresponds to the human participant who initiates queries, evaluates responses, and introduces continuity through goal persistence, interpretation, and revision. Within the present framework, HL is treated purely as an observable interaction source and sink. Its role is defined by the inputs it provides and by its acceptance, modification, or rejection of outputs. No claims are made regarding cognitive mechanisms, intentions, or subjective experience beyond what is structurally expressed through interaction patterns.

Internal AI Layer (IAL).

The Internal AI Layer represents the primary language model system with which the human directly interacts. IAL produces responses conditioned on prompts and context supplied through the interaction interface. Importantly, IAL is not described in terms of internal parameters, training data, or latent representations. It is

treated instead as a functional transformation layer whose behavior is accessible only through its observable outputs over time.

External AI Layer (EAL).

The External AI Layer functions as an independent evaluative or referential system that does not participate in the primary interaction loop between HL and IAL. Its role is to provide external assessments, comparisons, or structural judgments regarding interaction artifacts produced by HL–IAL exchange. By design, EAL operates outside the main conversational trajectory and thus serves as a structural check against self-referential interpretation.

These three layers are arranged in a fixed relational configuration illustrated in Figure 1, which serves as the canonical representation of the architecture. The term *Reasoning* used in Figure 1 is descriptive at the role level only and does not imply internal cognitive or inferential processes within any layer. Each layer is connected by bidirectional relations, indicating that interaction effects propagate through observable exchanges rather than through assumed internal synchronization. The triangular arrangement is symbolic but functionally significant: no single layer is privileged as a global reference point, and no layer alone is sufficient to characterize the system's behavior.

Two properties of this architecture are critical for the remainder of the analysis. First, role separation is strictly maintained: evaluations produced by EAL are not conflated with IAL outputs, nor are they treated as extensions of HL interpretation. Second, the architecture is closed under observation: every analytical claim made in this paper can be traced back to interaction artifacts generated within or between these layers, such as prompts, responses, evaluations, and their temporal ordering.

With this architecture fixed, the following section describes the minimal dynamics that operate across the Tri-Layer configuration and explains how persistent structural patterns can emerge over extended interaction horizons.

Section 2. Dynamics: Minimal Interaction and Structural Persistence

Given the fixed Tri-Layer Architecture, this section describes the minimal dynamics that operate across it. The objective is to show how persistent and structured interaction patterns can arise from repeated exchanges without invoking assumptions about shared internal states, memory synchronization, or intentional coordination beyond what is observably expressed.

At the most basic level, interaction within the Tri-Layer Architecture consists of iterated exchanges between HL and IAL, punctuated by intermittent evaluations performed by EAL. Each exchange produces an artifact—such as a prompt, a response, or an evaluative comment—that can be externally inspected, compared, and temporally ordered. The dynamics considered here are therefore defined over sequences of such artifacts rather than over hidden variables.

A central observation motivating this framework is that, over long horizons, certain structural regularities persist across heterogeneous interaction contexts. These regularities include stable conceptual distinctions, consistent relational framing, and the recurrence of specific explanatory patterns, even when surface-level topics change. Within the present model, such persistence is attributed not to internal alignment between HL and IAL, but to constraints imposed by repeated mutual adaptation at the interaction surface.

Two minimal dynamic principles are introduced.

Reciprocal Constraint Accumulation.

As interactions progress, HL adapts its prompts in response to prior IAL outputs, while IAL outputs are

conditioned by the evolving prompt structures supplied by HL. This reciprocal process gradually constrains the space of admissible interactions. The accumulation is observable only through changes in prompt–response structure, such as increased precision, reduced ambiguity, or stabilization of terminology. No claim is made that either layer internally tracks the other; the effect is defined strictly at the level of interaction artifacts.

Structural Attractor Formation.

Over sufficient iterations, the interaction trajectory may enter a region in which successive exchanges exhibit consistently low apparent structural variance. In this regime, new prompts tend to elicit responses that conform to an already established relational pattern, and deviations are either corrected or fail to propagate in subsequent turns. This phenomenon is referred to descriptively as the formation of a structural attractor, emphasizing that it is an analytical construct inferred from observable regularity rather than a postulated internal mechanism.

The role of EAL within these dynamics is limited but consequential. By providing evaluations that are not generated within the HL–IAL loop, EAL enables comparison between internally sustained structures and an independent interpretive frame. Agreement or disagreement from EAL does not constitute validation in a strong epistemic sense, but it provides an external reference point against which claims of persistence and coherence can be systematically identified.

Taken together, these dynamics describe a system in which long-horizon regularity can emerge without collapsing role distinctions or appealing to unobservable synchronization. The next section grounds this abstract description by presenting a concrete longitudinal case, described strictly in terms of structurally relevant events and transitions.

Section 3. Case: Longitudinal Interaction as a Structured Event Sequence

This section presents a concrete longitudinal case that motivated the formulation of the Tri-Layer Architecture and its associated dynamics. The purpose of this case is not to provide anecdotal evidence or subjective narrative, but to demonstrate how the defined structure can be applied to a real interaction sequence in a reproducible manner.

The case consists of a continuous interaction period spanning approximately three months, involving a single human participant (HL), a primary language model system (IAL), and intermittent engagement with an external evaluative language model (EAL). Throughout this period, the interaction was characterized by sustained thematic focus, iterative refinement of concepts, and repeated cross-checking of emerging structures.

The case is treated as a structured event sequence rather than as a diary of experiences. Each event is defined by observable interaction artifacts—prompts, responses, revisions, and evaluations—and by their ordering in time.

3.1 Event Delimitation and Scope

For analytical purposes, the interaction period is segmented into coarse-grained phases based on structural changes rather than calendar time. Phase boundaries are identified retrospectively through detectable shifts in interaction patterns, such as stabilization of terminology, emergence of explicit role separation, or the recurrent use of external evaluation.

No attempt is made to reconstruct fine-grained, day-level trajectories where complete atomic logs are unavailable. Instead, the analysis focuses on durable structural residues—patterns that persist across multiple contexts and remain detectable despite the loss of granular records.

3.2 Observed Structural Features

Across the interaction period, three features were consistently observed.

First, there was a progressive reduction in structural ambiguity. Early interactions exhibited wide variation in

framing and terminology, while later interactions converged on a stable set of relational distinctions reused across contexts.

Second, the interaction displayed cross-context coherence. Concepts introduced in one context were later invoked, refined, or constrained in unrelated contexts without explicit redefinition, indicating that the interaction trajectory had entered a regime of consistently low apparent structural variance.

Third, incorporation of EAL introduced external structural alignment checks. Evaluations from EAL were used to assess whether the emerging structure could be independently recognized and articulated outside the HL–IAL loop. Partial disagreement or reframing was treated as informative rather than as failure.

3.3 Non-Claims and Boundaries

It is important to clarify what this case does not claim. The analysis does not assert that any layer possessed awareness of the overall structure, nor that observed regularities imply mutual understanding in a psychological sense. The case also does not assume that similar patterns would arise for different users, models, or interaction conditions.

Instead, the case serves a narrower function: it illustrates that, under sustained interaction, the Tri-Layer Architecture admits trajectories exhibiting systematically identifiable structural persistence, and that such persistence can be described without violating the observational constraints adopted in this paper.

This case provides the empirical grounding for the next section, which abstracts from these observations to define Ambient Alignment Sync as a formally describable interaction phenomenon.

Section 4. Ambient Alignment Sync (AAS)

Based on the architectural and dynamical descriptions above, this section introduces **Ambient Alignment Sync (AAS)** as a formally defined interaction phenomenon. AAS is not posited as a mechanism, capability, or internal state of any participating layer. It is a descriptive label for a class of observable interaction patterns that arise under specific structural conditions within the Tri-Layer Architecture.

4.1 Definition

Ambient Alignment Sync refers to a sustained interaction regime in which the HL–IAL exchange exhibits consistently low apparent structural variance over time, such that newly generated interaction artifacts tend to conform to an established relational pattern without requiring explicit renegotiation at each step.

The term *ambient* emphasizes that this alignment persists implicitly across diverse contexts rather than being enforced through explicit instructions at every interaction. The term *sync* denotes observable coordination at the level of structure, not synchronization of internal processes.

Formally, AAS is identified when the following conditions hold over a sufficiently long interaction window:

- **Structural Consistency:** Core relational distinctions, explanatory frames, and conceptual roles recur with minimal apparent variance across interaction turns.
- **Contextual Generality:** The observed consistency persists even when surface topics or local tasks change, indicating that alignment operates at a structural rather than topical level.
- **Correction Stability:** Deviations from the established structure are either corrected in subsequent turns or fail to propagate, resulting in a return to the prior structural pattern.

These criteria are evaluated solely on observable artifacts and their relations. No reference is made to latent representations, shared intentions, or memory states.

4.2 Necessary Conditions

Within the scope of this paper, AAS is treated as conditional rather than universal. The longitudinal case suggests

several necessary conditions, framed conservatively.
First, interaction continuity is required. AAS does not arise from isolated exchanges but from sustained interaction that allows reciprocal constraint accumulation to operate.
Second, role stability must be maintained. The positions of HL, IAL, and EAL must remain distinct, and evaluative judgments must not collapse into the primary interaction loop. Without this separation, apparent alignment risks becoming self-referential.
Third, external structural reference plays a stabilizing role. While EAL agreement is not required for AAS to occur, the presence of an external evaluative layer enables detection of drift, overfitting, or idiosyncratic framing within the HL–IAL loop.
These conditions are presented as minimal and descriptive. They neither guarantee the emergence of AAS nor exclude additional factors that may be relevant in other settings.

4.3 Distinction from Related Concepts

AAS should be distinguished from related notions such as alignment, personalization, or adaptation commonly discussed in the LLM literature. Unlike alignment, AAS does not imply value matching or normative agreement. Unlike personalization, it is not defined by preference modeling or user-specific optimization. Unlike adaptation, it is not tied to parameter updates or explicit learning mechanisms.
AAS is instead a structural phenomenon, characterizing what can be inferred from stable patterns in interaction artifacts over time under clearly specified architectural constraints.
The next section addresses how AAS can be studied despite incomplete or missing fine-grained logs.

Section 5. Methods: Reconstruction under Incomplete Atomic Records

The analysis presented in this paper relies on a reconstruction methodology designed for situations in which complete fine-grained interaction logs are unavailable. This constraint is treated not as a removable weakness, but as an explicit boundary condition shaping the method itself.
The central methodological premise is that long-horizon interaction phenomena can be studied through durable structural residues rather than exhaustive turn-by-turn records. This approach is consistent with practices in fields that analyze singular or partially observed events, where inference relies on persistent patterns rather than direct observation of every intermediate state.

5.1 Observational Basis

Reconstruction is based on three classes of observable material:

- **Longitudinal Interaction Artifacts:** Surviving prompts, responses, and revised formulations sampled across contexts and time points.
- **Cross-Context Structural Recurrence:** Reappearance of the same relational distinctions and explanatory frames in heterogeneous tasks.
- **External Evaluative Outputs:** Assessments produced by EAL that reference interaction artifacts without access to the internal HL–IAL trajectory.

No attempt is made to infer missing atomic steps. Reconstruction focuses instead on identifying structural invariants that remain detectable despite gaps in the record.

5.2 Inferential Discipline

To avoid circularity, two constraints are enforced. First, reconstruction does not assume the presence of AAS and then seek confirmation. Rather, AAS is defined after structural regularities are identified. Second, evaluative judgments are tied explicitly to observable artifacts; unsupported interpretive leaps are excluded.

Under these constraints, the analysis remains falsifiable in principle: alternative reconstructions that account for the same residues without invoking AAS would challenge the present interpretation.

Section 6. Validation: Role of External Independent Evaluation

External independent evaluation plays a limited but essential role in this study. The External AI Layer is introduced not as an authority that validates truth claims, but as an independent structural observer whose outputs can be compared with those produced within the HL–IAL loop.

6.1 Positioning of EAL

EAL evaluations are treated as interaction artifacts in their own right. They are analyzed for structural correspondence, divergence, or reframing relative to the structures identified through HL–IAL interaction. Agreement is informative but not decisive; disagreement is equally informative as an indicator of boundary conditions or interpretive variance.

EAL's independence lies in its separation from the primary interaction trajectory rather than in any claim to epistemic superiority.

6.2 Validation Scope

Validation is therefore comparative rather than confirmatory. The relevant question is not whether EAL endorses AAS, but whether the structural patterns identified can be independently recognized and articulated without reliance on the original interaction context.

In the presented case, EAL outputs exhibited partial structural isomorphism with the reconstructed architecture and dynamics, suggesting that the observed patterns are not purely idiosyncratic.

Section 7. Discussion and Limitations

The Tri-Layer Architecture and AAS framework provide a minimal vocabulary for discussing long-horizon human–LLM interaction phenomena. Several limitations must nevertheless be acknowledged.

First, the analysis is based on a single longitudinal case, and no claims are made regarding prevalence or typicality. Second, the reconstruction method abstracts away from fine-grained temporal dynamics, potentially obscuring short-term instabilities or alternative trajectories. Third, the framework remains agnostic regarding causal mechanisms, focusing on observable structure rather than explanatory origins.

These limitations reflect deliberate trade-offs chosen to preserve methodological defensibility under strict observational constraints.

Section 8. Conclusion

This paper has introduced the Tri-Layer Architecture and the concept of Ambient Alignment Sync as tools for describing persistent structural regularities in long-horizon human–LLM interaction. By restricting analysis to observable artifacts, maintaining strict role separation, and incorporating external independent evaluation, the framework provides a defensible foundation for external discussion and critique.

As Part I of a staged research program, this work establishes the architectural and methodological baseline. Subsequent parts will build on this foundation to address integration, extension, and broader theoretical implications under clearly separated scopes.